

Introduction to Data Management

DSC 100

Nested SQL Queries

What have we learned so far

SQL features

- Projections
- Selections
- Joins (inner and outer)
- Aggregates
- Group by
- Having
- Inserts, updates, and deletes

Subqueries

- **Subquery:** A query that is part of another
- **Nested Query:** A query that has an embedded subquery
- A subquery can be a nested query itself!
- Why? Sometimes we need to express a condition that refers to a table that must itself be computed
- A subquery may occur in:
 - A SELECT clause
 - A FROM clause
 - A WHERE clause
- **Rule of thumb: avoid nested queries when possible**
 - But sometimes it's impossible, as we will see

Often appear here

Subqueries

- Can return a single value to be included in a **SELECT** clause
- Can return a relation to be included in the **FROM** clause
- Can return a single value to be compared with another value in a **WHERE** or **HAVING** clause
- Can return a relation to be used in the **WHERE**

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actortname, moviename, salary)`

1. Subqueries in SELECT

For each actor return the genre of movie they acted in

```
SELECT a.actortname, (SELECT genre  
                      FROM Movie m  
                     WHERE m.name=a.moviename) as genre  
FROM ActedIn a
```

“correlated
subquery”

What happens if the subquery returns more than one genre?

```
Movie(name, year, genre, budget, revenue , rating)
```

```
ActedIN(actortname, moviename, salary)
```

1. Subqueries in SELECT

Whenever possible, don't use a nested queries:

```
SELECT a.actortname, (SELECT genre  
                      FROM Movie m  
                     WHERE m.name=a.moviename) as genre  
FROM ActedIn a
```



Subquery unnesting

```
SELECT a.actortname, genre  
FROM ActedIN a, Movie m  
WHERE m.name=a.moviename
```

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actortname, moviename, salary)`

1. Subqueries in SELECT

Compute average salary of actors for all movies with rating > 9

```
SELECT DISTINCT m.name, (SELECT AVG(salary)
                           FROM   ActedIn a
                           WHERE  m.name=a.moviename) as salary
FROM Movie m
WHERE m.rating > 9
```

Subquery
unnesting

```
SELECT m.name, AVG(salary)
FROM   Movie m, ActedIn a
WHERE  m.name=a.moviename
AND   m.rating > 9
GROUP BY m.name
```

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actortname, moviename, salary)`

1. Subqueries in SELECT

```
SELECT DISTINCT m.name, (SELECT count(*)
                           FROM ActedIN a
                           WHERE m.name=a.moviename) as anum
  FROM Movie m
```

Movie(name, year, genre, budget, revenue , rating)

ActedIN(actortname, moviename, salary)

1. Subqueries in SELECT

Compute the number of actors in each movie

```
SELECT DISTINCT m.name, (SELECT count(*)  
                           FROM ActedIN a  
                          WHERE m.name=a.moviename) as anum  
FROM Movie m
```

||



Subquery
unnesting

```
SELECT m.name, count(*)  
FROM   ActedIn a, Movie m  
WHERE  m.name=a.moviename  
GROUP BY m.name
```

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actortname, moviename, salary)`

1. Subqueries in SELECT

But are these equivalent?

```
SELECT DISTINCT m.name, (SELECT AVG(salary)
                           FROM ActedIn a
                           WHERE m.name=a.moviename) as salary
FROM Movie m
WHERE m.rating > 9
```

```
SELECT m.name, AVG(salary)
FROM Movie m, ActedIn a
WHERE m.name=a.moviename AND
m.rating > 9
GROUP BY m.name
```



```
SELECT m.name, AVG(salary)
FROM Movie m LEFT OUTER JOIN ActedIn a ON
m.name=a.moviename
WHERE m.rating > 9
GROUP BY m.name
```

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actornname, moviename, salary)`

2. Subqueries in FROM

Find all Movie with rating > 8 and < 9

```
SELECT x.name, rating  
FROM (SELECT *  
      FROM Movie AS m  
      WHERE rating > 8) as x  
WHERE x.rating < 9
```

“Not a correlated subquery”

A subquery whose result we called myTable

```
WITH myTable AS (SELECT * FROM Movie AS m WHERE rating > 8)  
SELECT x.name, x.rating  
FROM myTable as X  
WHERE x.rating < 9
```

Sub-query refactoring

DSC 100

11

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actortname, moviename, salary)`

2. Subqueries in FROM

Find all Movie with rating > 8 and < 9

```
SELECT x.name  
FROM (SELECT *  
      FROM Movie AS m  
      WHERE rating > 8) as x  
WHERE x.rating < 9
```



Subquery
unnesting

```
SELECT m.name, rating  
FROM   Movie m  
WHERE  m.rating < 9 AND m.rating > 8
```

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actortname, moviename, salary)`

3. Subqueries in WHERE

Find the name of actors who have acted in some Sci-Fi movie

Quantifier is a logical operator that specifies how many elements in the domain of discourse satisfy a property

Existential quantifiers

TRUE iff the subquery returns one or more records

"there exists", "there is at least one", or "for some"

Using **EXISTS**:

```
SELECT DISTINCT a.actortname  
FROM ActedIn a  
WHERE EXISTS (SELECT m.name  
              FROM Movie m  
              WHERE m.name=a.moviename AND  
                    m.genre='Sci-Fi' )
```

Movie(name, year, genre, budget, revenue , rating)

ActedIN(actortname, moviename, salary)

3. Subqueries in WHERE

Find the name of actors who have acted in some Sci-Fi movie

Existential quantifiers

Using IN:

Allow us to test
set membership

```
SELECT DISTINCT a.actortname
FROM ActedIn a
WHERE a.moviename IN (SELECT m.name
                      FROM Movie m
                      WHERE m.name=a.moviename AND
                            m.genre='Sci-Fi' )
```

Movie(name, year, genre, budget, revenue , rating)

ActedIN(actortname, moviename, salary)

3. Subqueries in WHERE

Find the name of actors who have acted in some Sci-Fi movie

Existential quantifiers

```
SELECT DISTINCT a.actortname  
FROM ActedIn a  
WHERE a.moviename IN (SELECT m.name  
                      FROM Movie m  
                      WHERE m.name=a.moviename AND  
                            m.genre='Sci-Fi' )
```

Subquery
unnesting

```
SELECT DISTINCT a.actortname  
FROM Movie m, ActedIn a  
WHERE m.name=a.moviename AND m.genre='Sci-Fi'
```

Movie(name, year, genre, budget, revenue , rating)

ActedIN(actortname, moviename, salary)

3. Subqueries in WHERE

Find the name of actors who have acted in some non-Sci-Fi movie

Existential quantifiers

```
SELECT DISTINCT a.actortname  
FROM ActedIn a  
WHERE a.moviename NOT IN (SELECT m.name  
                           FROM Movie m  
                           WHERE m.name=a.moviename AND  
                                 m.genre='Sci-Fi')
```

Subquery
unnesting

```
SELECT DISTINCT a.actortname  
FROM Movie m, ActedIn a  
WHERE m.name=a.moviename AND m.genre ≠ 'Sci-Fi'
```

Existential quantifiers are straightforward! 😊

Join queries essentially evaluate the presence of existential quantifiers



Universal quantifiers pose a challenge 😞

The SQL constructs that we have covered thus far do not have the capability to handle universal quantifiers.

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actornname, moviename, salary)`

3. Subqueries in WHERE

Retrieve all actor names that only acted on action movies

Same as: Every movie they acted on was an action movie

Universal quantifiers

“given any”, “for all”, “for every”

Step 1: Find all actor names that acted on some non-action movie

```
SELECT a.actornname  
FROM Movie m, ActedIn a  
WHERE m.name=a.moviename AND m.genre ≠ 'Action'
```

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actornname, moviename, salary)`

3. Subqueries in WHERE

Retrieve all actor names that only acted on action movies

Same as: every movies they acted on were action

Universal quantifiers

Step 1: Find all actor names that acted on some non-action movie

Step 2: Retrieve all the others (i.e., those that do not satisfy the result of Step 1)

```
SELECT a.actornname  
FROM Movie m, ActedIn a  
WHERE m.name=a.moviename AND m.genre ≠ 'Action'
```

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actornname, moviename, salary)`

3. Subqueries in WHERE

Retrieve all actor names that only acted on action movies

Same as: every movies they acted on were action

Universal quantifiers

Step 1: Find all actor names that acted on some non-action movie

Step 2: Retrieve all the others (i.e., those that do not satisfy the result of Step 1)

```
SELECT a.name  
FROM ActedIn a  
WHERE a.actornname NOT IN ( SELECT a.actornname  
                            FROM Movie m, ActedIn a  
                            WHERE m.name=a.moviename AND m.genre ≠ 'Action' )
```

`Movie(name, year, genre, budget, revenue , rating)`

`ActedIN(actorname, moviename, salary)`

3. Subqueries in WHERE

Retrieve all actor names that acted on at most two action movies

```
SELECT DISTINCT a.actorname
FROM ActedIn a
WHERE 2 >= (SELECT count(*)
             FROM Movie m, ActedIN b
             WHERE m.name=b.moviename AND
                   b.actorname=a.actorname AND
                   m.genre='Action' )
```

What does this query do?

```
SELECT DISTINCT a.actorname
FROM ActedIn a
WHERE 0 > (SELECT count(*)
             FROM Movie m, ActedIN b
             WHERE m.name=b.moviename AND
                   b.actorname=a.actorname AND
                   m.genre='Action' )
```

Movie(name, year, genre, budget, revenue , rating)

ActedIN(actortname, moviename, salary)

3. Subqueries in WHERE

What does this query do?

```
SELECT a.actortname
FROM ActedIn a
JOIN Movie m ON a.moviename = m.name
WHERE m.genre = 'Action'
GROUP BY a.actortname
HAVING count(*) >= 2
```

`Movie(name, year, genre, budget, revenue , rating)`
`ActedIN(actortname, moviename, salary)`

3. Subqueries in WHERE

Find all movie s.t. all their actors' salaries > \$100K

Universal quantifiers

```
SELECT m.name  
FROM movie m  
WHERE $100K < ALL (SELECT a.salary  
                      FROM ActedIn a  
                     WHERE m.name=a.moviename)
```

Not supported
in sqlite

`Movie(name, year, genre, budget, revenue , rating)`
`ActedIN(actortname, moviename, salary)`

3. Subqueries in WHERE

Find all movie s.t. all their actors' salaries > \$100K

Universal quantifiers

```
SELECT m.name  
FROM movie m  
WHERE $100K < ALL (SELECT a.salary  
                      FROM ActedIn a  
                     WHERE m.name=a.moviename)
```

Not supported
in sqlite

`Movie(name, year, genre, budget, revenue , rating)`
`ActedIN(actortname, moviename, salary)`

3. Subqueries in WHERE

Find all movie s.t. all their actors' salaries > \$100K

Universal quantifiers

```
SELECT m.name
FROM movie m
WHERE NOT EXISTS
(SELECT 1 FROM ActedIn a WHERE m.name = a.moviename
AND a.salary <= $100k)
```

Unnesting Universal Quantifiers

- Is it possible to unnest the *universal quantifier* query?

Definition: A query Q is **monotone** if:

- Whenever we add tuples to one or more input tables, the answer to the query will not lose any output tuple

```
SELECT a.actorname  
FROM   Movie m, ActedIn a  
WHERE  m.name=a.moviename AND m.genre='Crime'
```

Is this monotone?

Movie(name, year, genre)

ActedIN(actortname, moviename)

Monotone Queries

name	year	genre
Apocalypse now	1979	War
The god father	1972	Crime
Planet Earth II	2016	Nature documentary
Jack and Jill	2011	Comedy

actortname	moviename
Marlon Brando	Apocalypse now
Al Pacino	The god father
Marlon Brando	The god father

```
SELECT a.actortname  
FROM Movie m, ActedIn a  
WHERE m.name=a.moviename AND m.genre='Crime'
```

actortname
Marlon Brando
Al Pacino

Movie(name, year, genre)

ActedIN(actortname, moviename)

Monotone Queries

name	year	genre
Apocalypse now	1979	War
The god father	1972	Crime
Planet Earth II	2016	Nature documentary
Jack and Jill	2011	Comedy

actortname	moviename
Marlon Brando	Apocalypse now
Al Pacino	The god father
Marlon Brando	The god father
Al Pacino	Jack and Jill

```
SELECT a.actortname
FROM Movie m, ActedIn a
WHERE m.name=a.moviename AND m.genre='Crime'
```

Monotone

actortname
Marlon Brando
Al Pacino

Movie(name, year, genre)

ActedIN(actortname, moviename)

Monotone Queries

name	year	genre
Apocalypse now	1979	War
The god father	1972	Crime
Planet Earth II	2016	Nature documentary
Jack and Jill	2011	Comedy

actortname	moviename
Marlon Brando	Apocalypse now
Al Pacino	The god father
Marlon Brando	The god father

```

SELECT a.actortname
FROM ActedIn a
WHERE a.actortname
NOT IN ( SELECT a.actortname
          FROM Movie m, ActedIn a
          WHERE m.name=a.moviename AND m.genre ≠ 'Crime' )
    
```

actortname
Al Pacino

Movie(name, year, genre)
ActedIN(actortname, moviename)

Monotone Queries

name	year	genre
Apocalypse now	1979	War
The god father	1972	Crime
Planet Earth II	2016	Nature documentary
Jack and Jill	2011	Comedy

actortname	moviename
Marlon Brando	Apocalypse now
Al Pacino	The god father
Marlon Brando	The god father
Al Pacino	Jack and Jill

```
SELECT a.actortname
FROM ActedIn a
WHERE a.actortname
NOT IN ( SELECT a.actortname
          FROM Movie m, ActedIn a
          WHERE m.name=a.moviename AND m.genre ≠ 'Crime' )
```

actortname
Al Pacino

Movie(name, year, genre)
ActedIN(actortname, moviename)

Monotone Queries

name	year	genre
Apocalypse now	1979	War
The god father	1972	Crime
Planet Earth II	2016	Nature documentary
Jack and Jill	2011	Comedy

actortname	moviename
Marlon Brando	Apocalypse now
Al Pacino	The god father
Marlon Brando	The god father
Al Pacino	Jack and Jill

```
SELECT a.actortname
FROM ActedIn a
WHERE a.actortname
NOT IN ( SELECT a.actortname
          FROM Movie m, ActedIn a
          WHERE m.name=a.moviename AND m.genre ≠ 'Crime' )
```

actortname

Non-monotone

Monotone Queries

- Theorem: If Q is a SELECT-FROM-WHERE query that does not have subqueries, and no aggregates, then it is monotone.
- Proof. We use the nested loop semantics: if we insert a tuple in a relation R_i , this will not remove any tuples from the answer

```
SELECT a1, a2, ..., ak
FROM   R1 AS x1, R2 AS x2, ..., Rn AS xn
WHERE  Conditions
```

```
for x1 in R1 do
  for x2 in R2 do
    ...
  for xn in Rn do
    if Conditions
      output (a1,...,ak)
```

Add a tuple to $R_2 \dots$

...can't lose anything here.

Non-Monotone Queries

- If a query is **monotonic**, it implies that a nested query can be unnested.
- If a query is **non-monotone**, it implies that a nested query can NOT be unnested
- Queries with universal quantifiers or negation and aggregates are non-monotone